

Application for Software Engineer III, AI/ML, DSP/HTP Optimisation

Dear Google Recruitment Team,

Why Me?

I am applying for the Software Engineer III, AI/ML, DSP/HTP Optimisation role with 7+ years of production engineering experience, hands-on AI/ML delivery in Python, and a record of turning performance bottlenecks into measured improvements. At JPMorgan Chase, I cut messaging turnaround by approximately 85%, from 30 minutes to under 5 minutes at 150K+ messages per hour, by redesigning a Java service around Kafka and asynchronous processing. That is the engineering pattern I would bring to this role: profile the critical path, understand the system constraints, make targeted changes, and validate the result at realistic scale.

At EPAM, I built a Python RAG agent that saves investment committee members an average of 2 hours of preparation per meeting, and I lead a seven-engineer team responsible for a cloud data pipeline that reconciles 500K+ entities per day across 12 source databases. My Google Professional ML Engineer certification complements practical experience with model deployment, evaluation, data processing, MLOps, Vertex AI, BigQuery, TensorFlow/Keras, Docker, and Kubernetes. Outside client work, I customize Linux environments and run local models with llama.cpp, making compute, memory, and hardware constraints tangible rather than abstract.

Why Google?

The Platforms and Devices challenge is compelling because model quality alone is not enough: useful on-device AI must also meet strict latency, power, thermal, and stability limits. I want to work where profiling data drives hardware-specific optimization, and where software decisions can shape efficiency standards for HTP, DSP, and next-generation devices.

I also bring the versatility and leadership expected of a Google engineer. I have owned architecture through production, worked across backend, cloud, data, and ML systems, and mentored engineers while remaining hands-on. I would contribute a disciplined optimization mindset, production Python and ML infrastructure experience, and the systems breadth needed to help make advanced AI features efficient and dependable on real devices.

Sincerely,
Spas Zahariev

Attached: Curriculum Vitae